


Generalized linear models for cracking the neural code




Jonathan Pillow

Princeton Neuroscience Institute, Psychology,
Center for Statistics & Machine Learning
Princeton University

Data Science and Data Skills for Neuroscientists
SFN Short Course
Nov 11, 2016

Retinal responses to white noise

(ON parasol cells)





Shlens, Field, Gauthier, Greschner, Sher, Litke & Chichilnisky (2009).

neural coding problem



X

stimuli

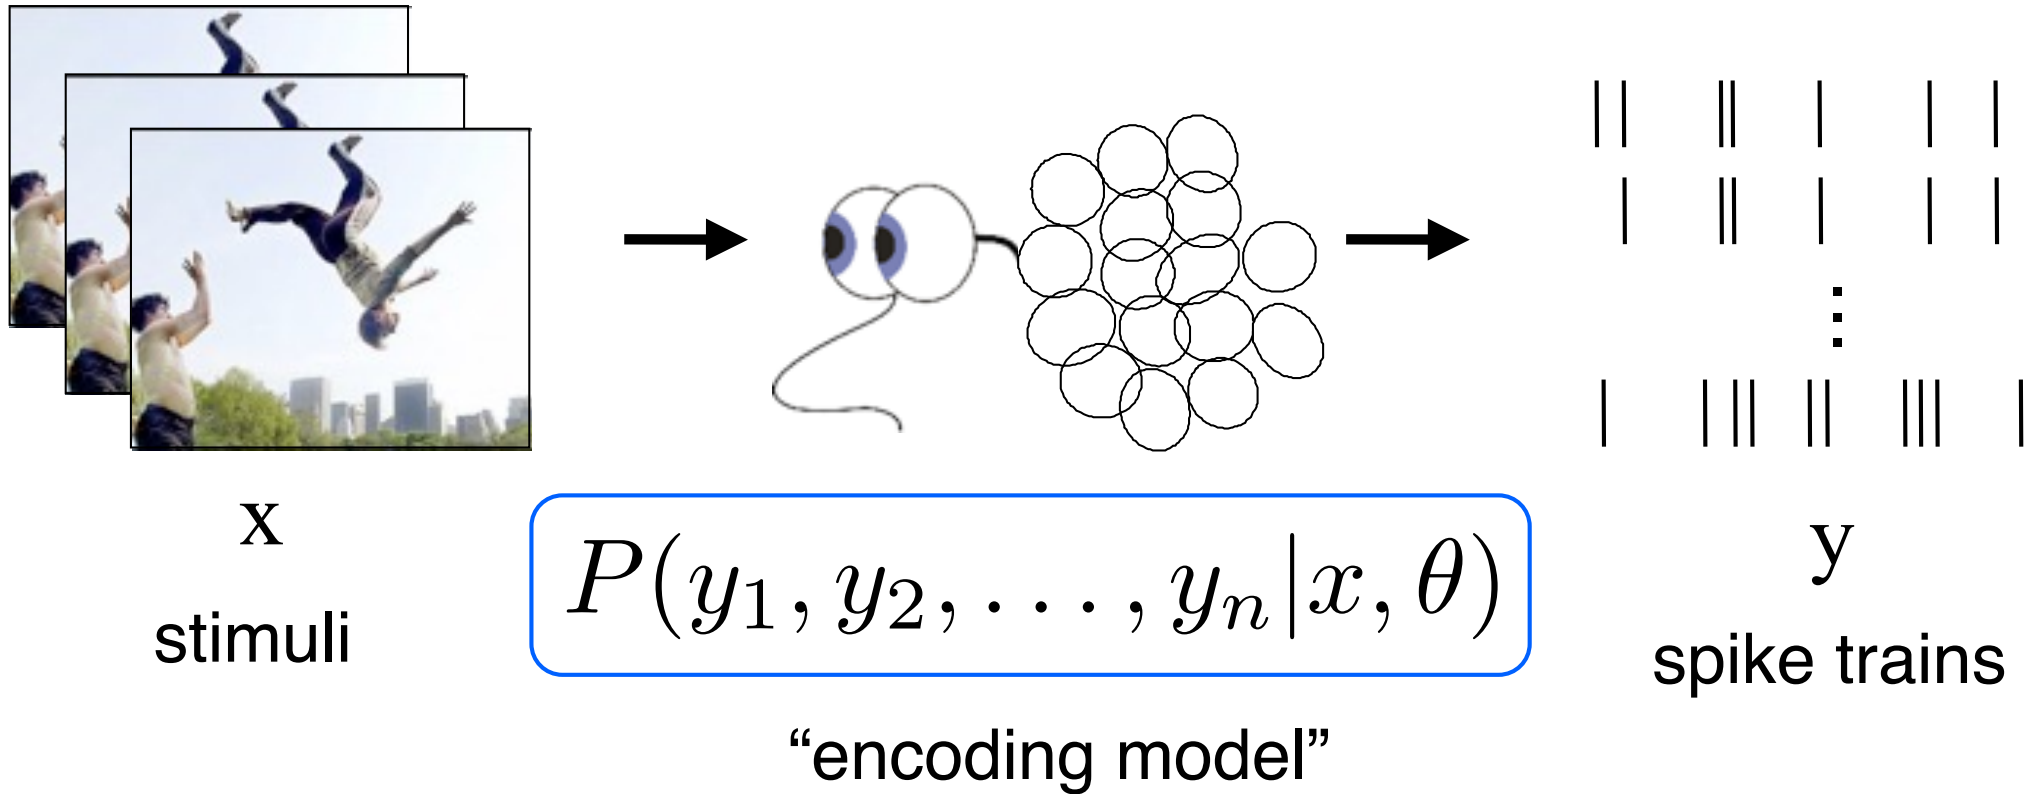


y

spike trains

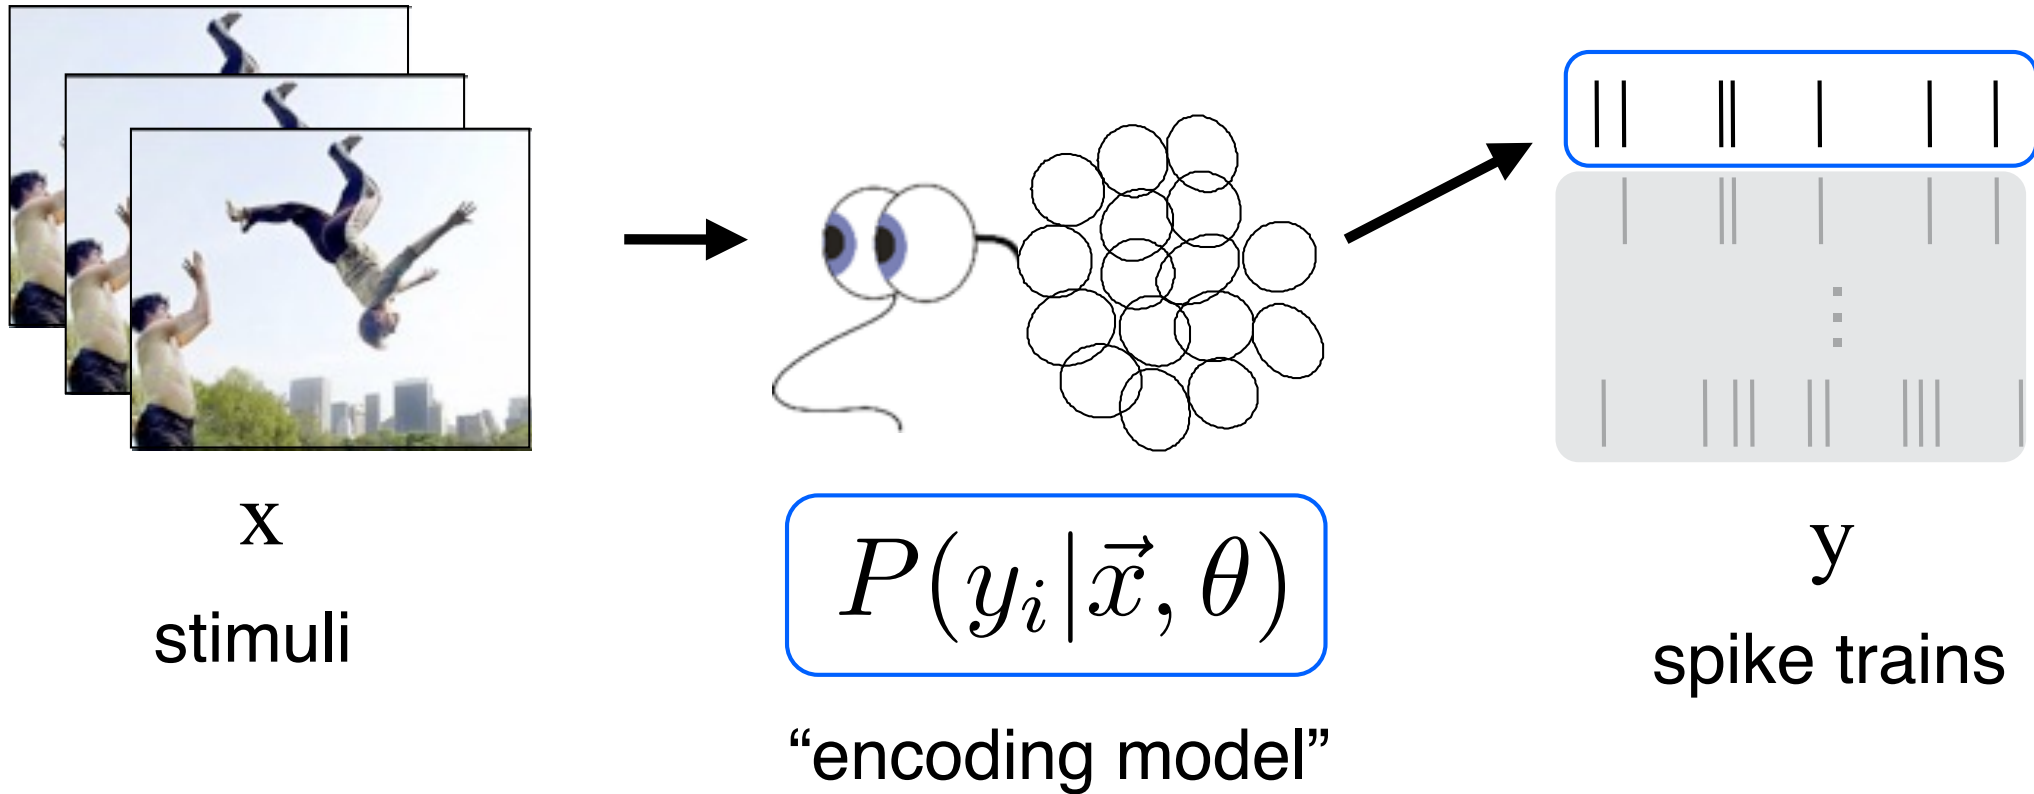
Q: what is the probabilistic relationship between stimuli and spike trains?

neural coding problem



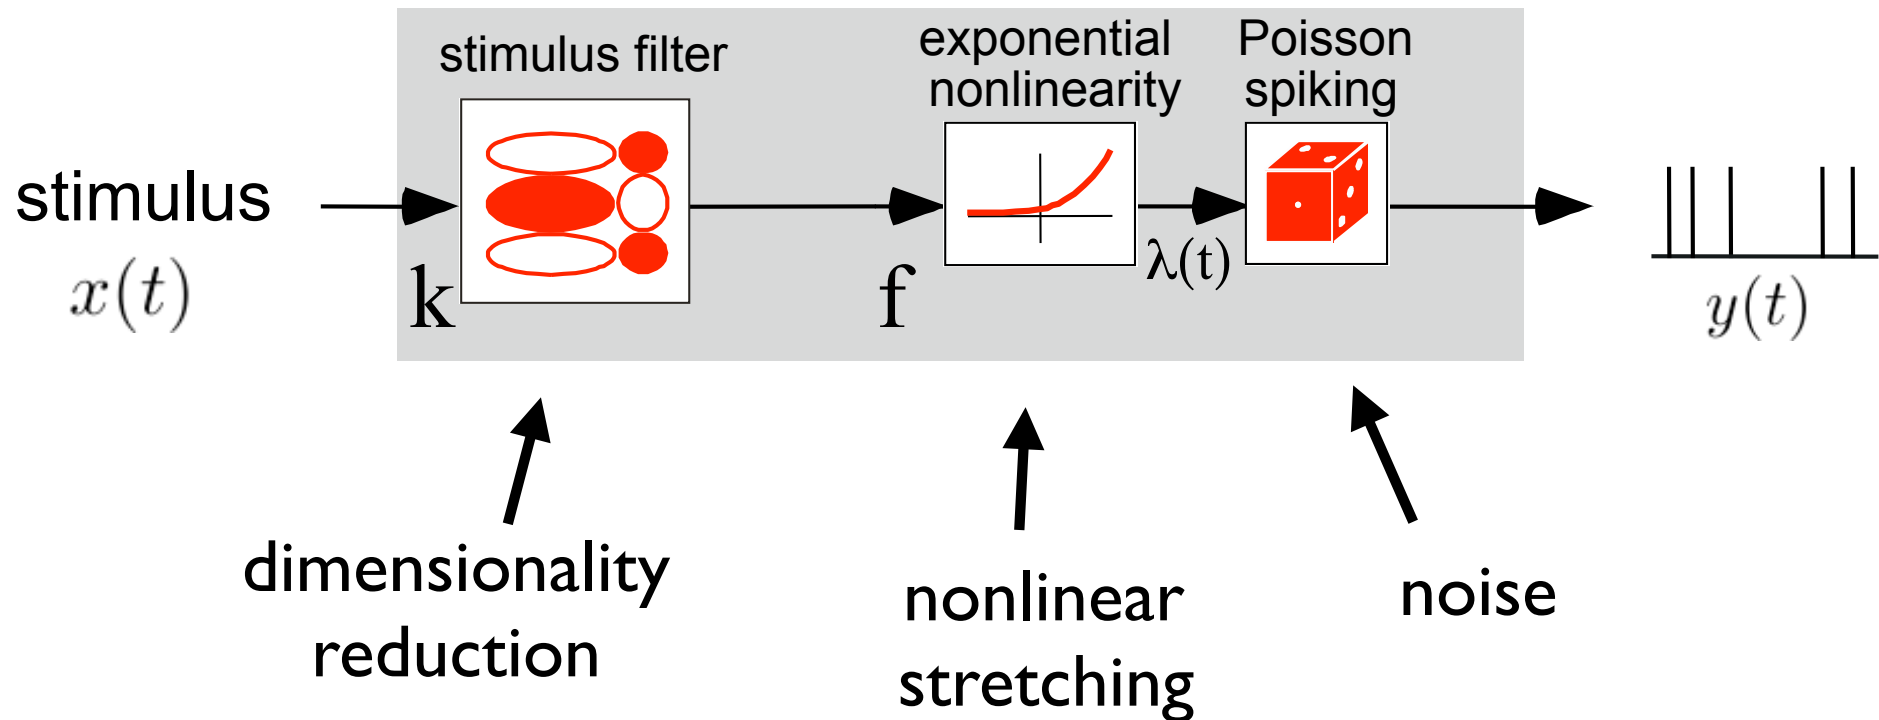
Q: what is the probabilistic relationship between stimuli and spike trains?

first block: single-neuron encoding



“basic” Poisson GLM

Linear-Nonlinear-Poisson model



spike rate $\lambda = f(\vec{k} \cdot \vec{x})$

spike count $y \sim \text{Pois}(\lambda)$

What is a GLM?

Be careful about terminology:

GLM

≠

GLM

General Linear Model

Generalized Linear Model

(Nelder 1972)

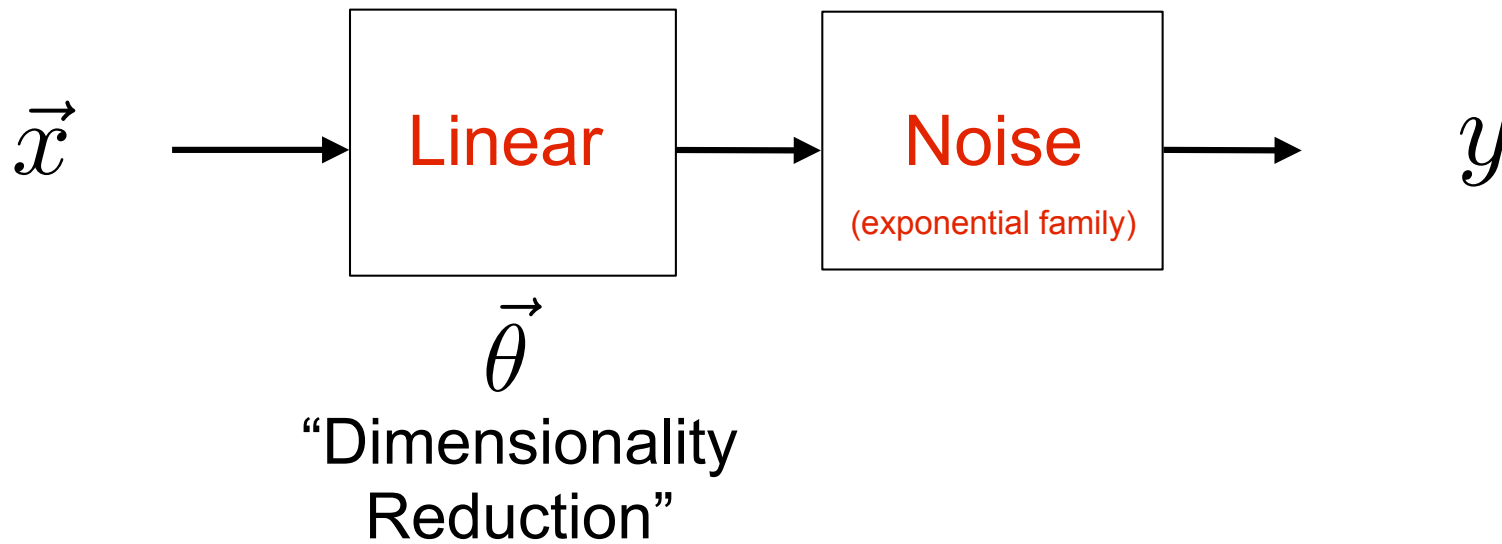
Linear

~~Linear~~

Moral:


Be careful when naming your model!

1. General Linear Model



- Examples:
1. Gaussian $y = \vec{\theta} \cdot \vec{x} + \epsilon$
 2. Poisson $y \sim \text{Pois}(\vec{\theta} \cdot \vec{x})$

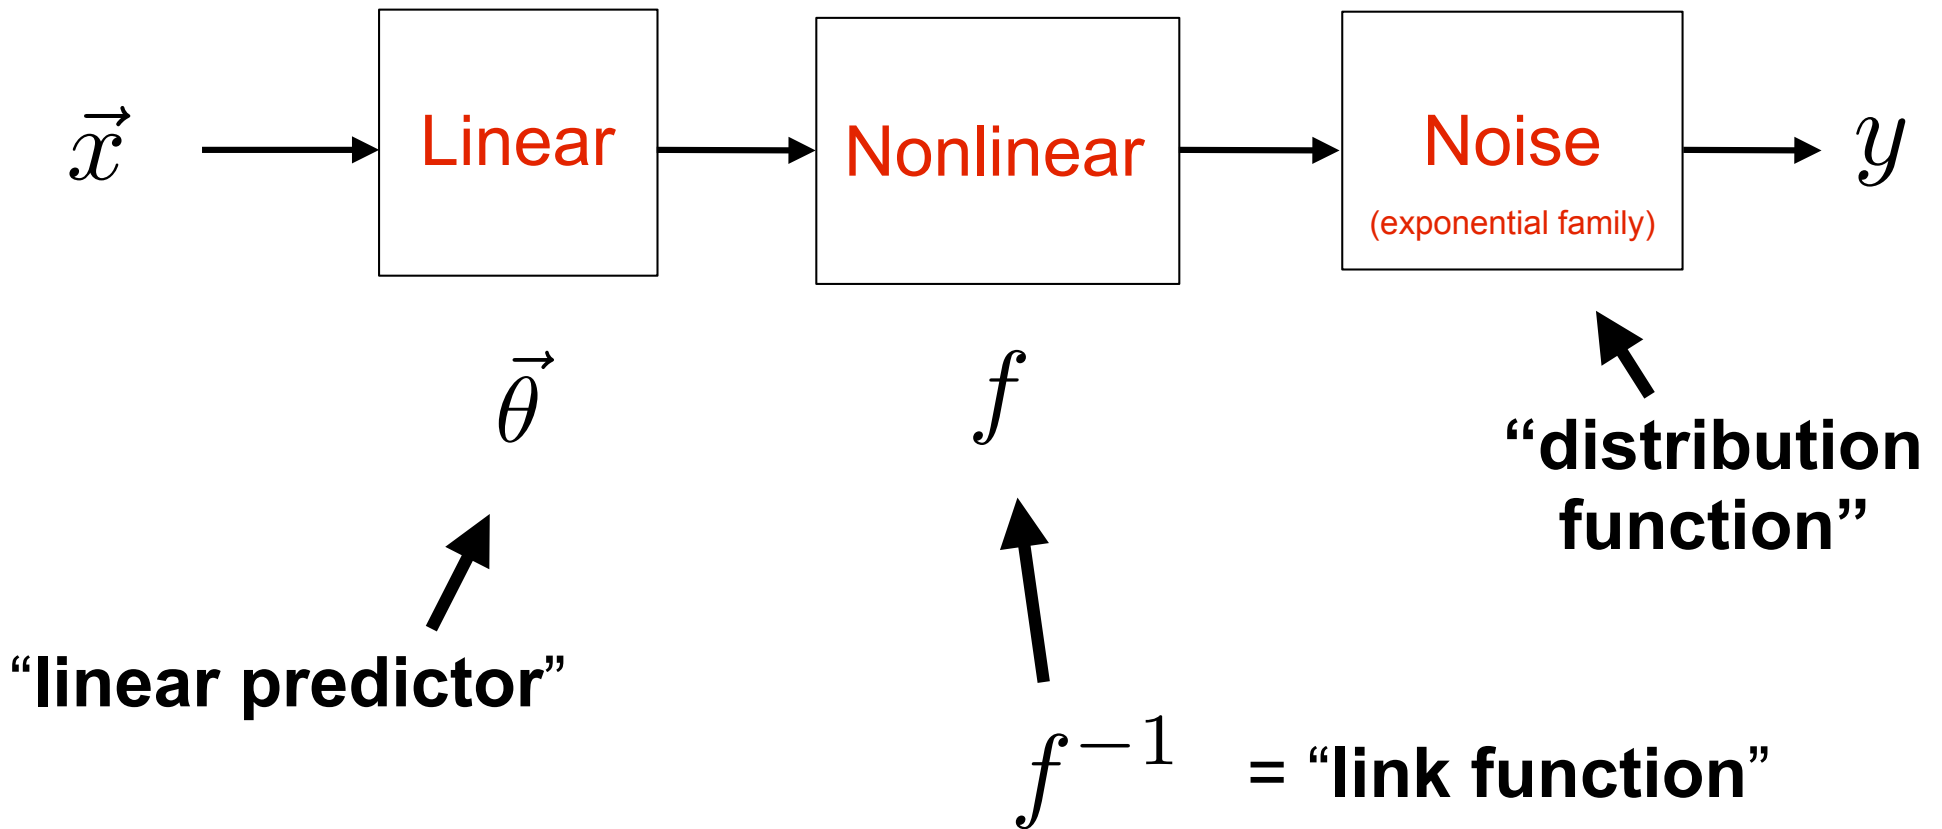
2. Generalized Linear Model



- Examples:
1. Gaussian $y = f(\vec{\theta} \cdot \vec{x}) + \epsilon$
 2. Poisson $y \sim \text{Pois}(f(\vec{\theta} \cdot \vec{x}))$

2. Generalized Linear Model

Terminology:



Applying it to data


response
at time t

$$y_t = \vec{k} \cdot \vec{x}_t + \text{noise}$$


linear
filter

vector stimulus
at time t

stimulus



response



time →

response
at time t

$$y_t = \vec{k} \cdot \vec{x}_t + \text{noise}$$

linear
filter

vector stimulus
at time t

linear filter

vector stimulus
at time t


walk through the data
one time bin at a time

$$\vec{x}_t$$

stimulus

response

time \longrightarrow

 y_t 

response
at time t

$$y_t = \vec{k} \cdot \vec{x}_t + \text{noise}$$

linear
filter

vector stimulus
at time t


walk through the data
one time bin at a time

$t = 2$

stimulus

\vec{x}_t

response



response
at time t

$$y_t = \vec{k} \cdot \vec{x}_t + \text{noise}$$

linear
filter

vector stimulus
at time t


walk through the data
one time bin at a time

$t = 3$

stimulus

\vec{x}_t

response



response
at time t

$$y_t = \vec{k} \cdot \vec{x}_t + \text{noise}$$

linear
filter

vector stimulus
at time t


walk through the data
one time bin at a time

$t = 4$

stimulus

\vec{x}_t

response



response
at time t

$$y_t = \vec{k} \cdot \vec{x}_t + \text{noise}$$

linear
filter

vector stimulus
at time t


walk through the data
one time bin at a time

$t = 5$

stimulus

\vec{x}_t

response



Build up to following matrix version:

$$\begin{array}{c} \text{time} \\ \downarrow \end{array} \quad Y = X \vec{k} + \text{noise}$$
$$\begin{bmatrix} 0 \\ 0 \\ 1 \\ \vdots \end{bmatrix} = \underbrace{\begin{bmatrix} \text{blue step function} \\ \text{blue step function} \\ \text{blue step function} \\ \vdots \end{bmatrix}}_{\text{design matrix}} \begin{bmatrix} \vec{k} \end{bmatrix}$$

The diagram illustrates the matrix formulation of a linear model. On the left, a vertical vector Y is shown with elements 0, 0, 1, and a vertical ellipsis. To its left, a vertical arrow labeled "time" points downwards. This vector is equal to the product of a matrix X and a vector \vec{k} , plus a noise term. The matrix X is represented by a large square bracket containing three blue step functions and a vertical ellipsis. A red curly brace underneath this matrix is labeled "design matrix". The vector \vec{k} is shown as a vertical bracket containing the symbol \vec{k} . Arrows point from the Y and X in the first equation to their respective matrix representations in the second equation.

Build up to following matrix version:

$$Y = X \vec{k} + \text{noise}$$

time ↓

$$\begin{bmatrix} 0 \\ 0 \\ 1 \\ \vdots \end{bmatrix} = \begin{bmatrix} \text{stimulus trace 1} \\ \text{stimulus trace 2} \\ \text{stimulus trace 3} \\ \vdots \end{bmatrix} \begin{bmatrix} \vec{k} \end{bmatrix}$$

I. “Linear-Gaussian” GLM: $\hat{k} = \underbrace{(X^T X)^{-1}}_{\text{stimulus covariance}} \underbrace{X^T Y}_{\text{spike-triggered avg (STA)}}$

Build up to following matrix version:

$$Y = X \vec{k} + \text{noise}$$

time ↓

$\begin{bmatrix} 0 \\ 0 \\ 1 \\ \vdots \end{bmatrix} = \begin{bmatrix} \text{step function} \\ \text{step function} \\ \text{step function} \\ \vdots \end{bmatrix} \begin{bmatrix} \vec{k} \end{bmatrix}$

1. “Linear-Gaussian” GLM: $\hat{k} = (X^T X)^{-1} X^T Y$

2. Poisson GLM: `k = glmfit(X, Y, 'Poisson');`

maximum likelihood fit (assumes exponential nonlinearity by default)

Now do it: tutorial1_PoissonGLM.m

1. cd into correct directory ('pillow_GLMtutorials')
2. open tutorial: use `ctrl-enter`, `ctrl-down`

Dataset:

- retinal ganglion cell spike trains [Uzzell & Chichilnisky 2004]
- 2 OFF cells, 2 on cells
- full-field, binary white noise stimuli

Now do it: tutorial1_PoissonGLM.m

1. cd into correct directory ('pillow_GLMtutorials')
2. open tutorial: use ctrl-enter, ctrl-down

Topics:

- 1-2. load data, bin spike trains
3. build design matrix (slow and fast versions)
- 4a. compute STA (for vizualization)
- 4b. linear-Gaussian GLM ("least-squares regression")
5. Poisson GLM: fit assuming exponential nonlinearity
6. Poisson GLM: non-parametric estimate of nonlinearity
7. model comparison: log-likelihood & mutual information
8. model comparison: AIC
9. simulate from fitted model

Contrary to what Konrad may have told you:
do **NOT** use 'fminsearch'!

DO use:

1. `fminunc` - unconstrained optimization
2. `fmincon` - constrained optimization



Block II:

GLMs with spike-history and coupling



X


stimuli



y

spike trains

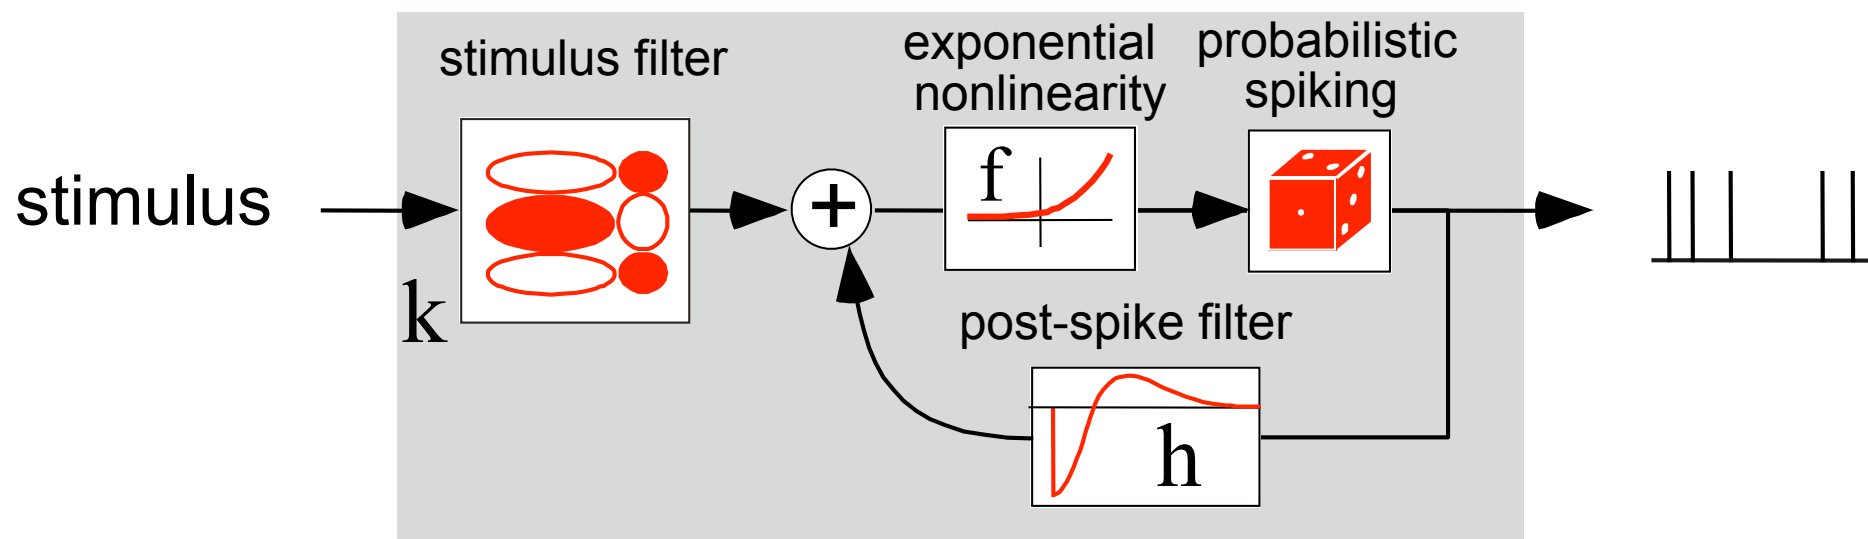
Poisson GLM



spike rate $\lambda(t) = f(k \cdot x(t))$

- problem: assumes spiking depends only on stimulus!

Poisson GLM with spike-history dependence




spike rate:
$$\lambda(t) = f(\vec{k} \cdot \vec{x}(t) + \vec{h} \cdot \vec{y}_{hst}(t))$$

$$= e^{\vec{k} \cdot \vec{x}(t)} \cdot e^{\vec{h} \cdot \vec{y}_{hst}(t)}$$

- output: no longer a Poisson process
- interpretation: “soft-threshold” integrate-and-fire model



GLM dynamic behaviors

- irregular spiking




GLM dynamic behaviors

- regular spiking




GLM dynamic behaviors

- adaptation




post-spike filter
 $h(t)$

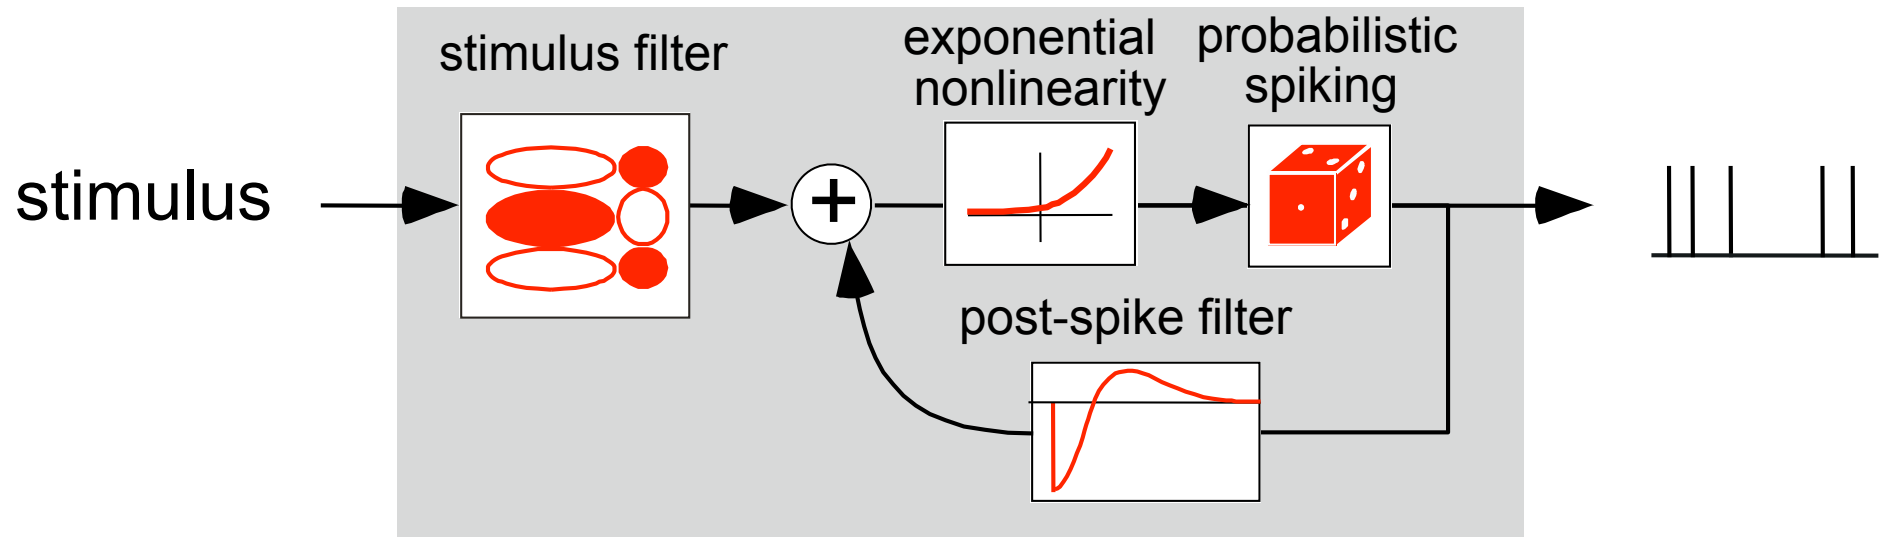


GLM dynamic behaviors


- bursting




Generalized Linear Model (GLM)




multi-neuron GLM



multi-neuron GLM



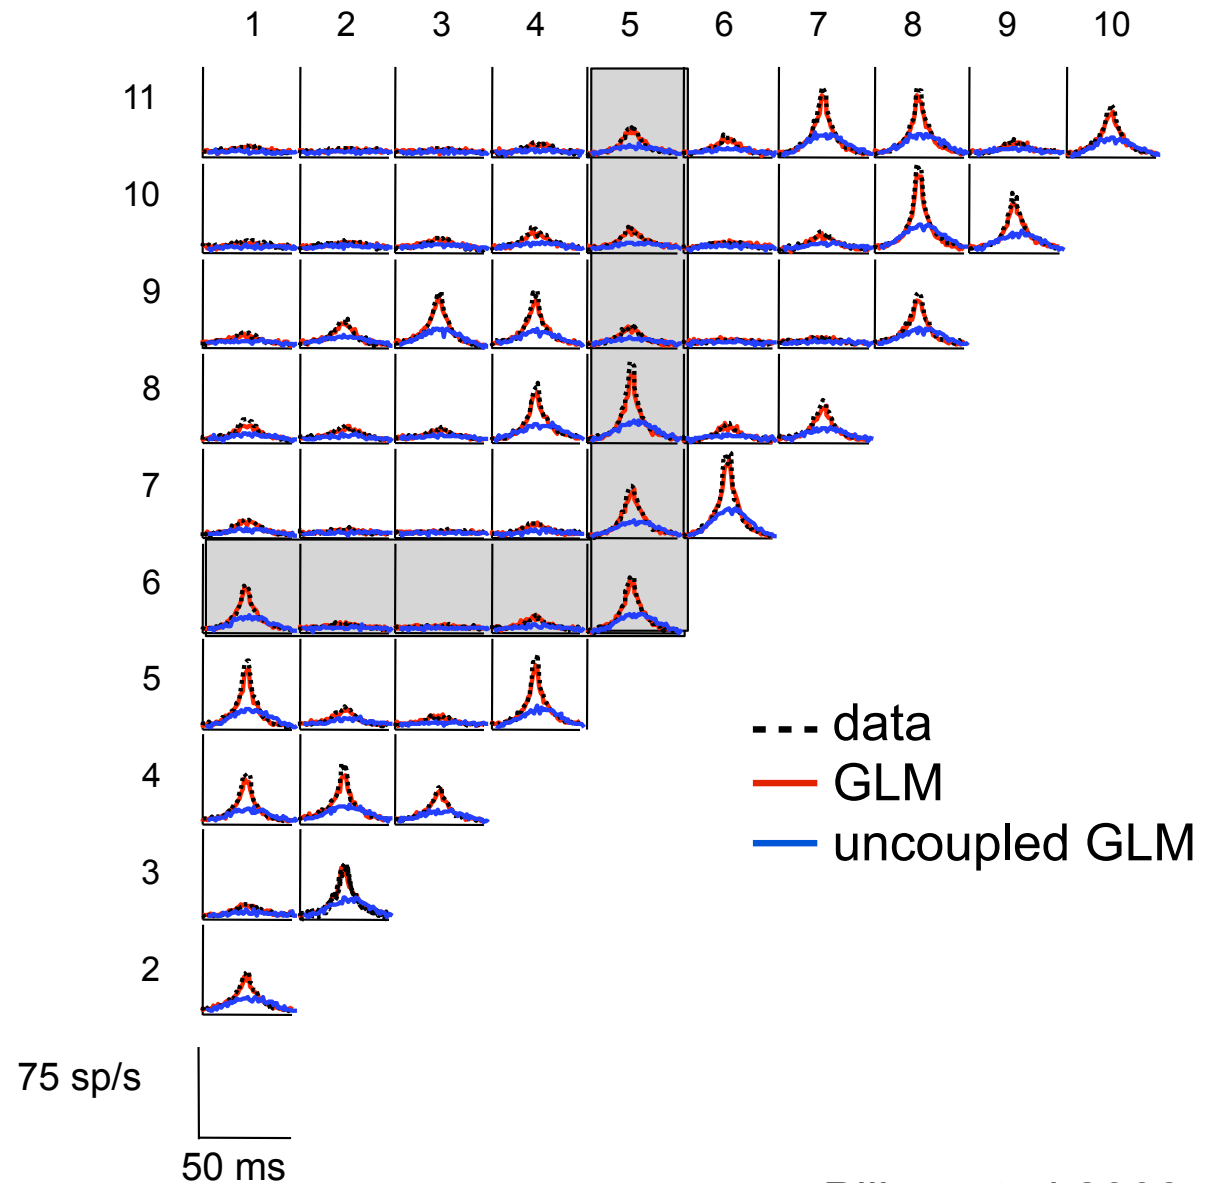

GLM equivalent diagram:



spike rate $\lambda_i(t) = \exp(k_i \cdot x(t) + \sum_j h_{ij} \cdot y_j(t))$


modeling correlation structure in neural spike trains

receptive fields




Fitting: Maximum Likelihood

Data




GLM



- find filters that maximize the log-conditional probability of the observed data

$$\log P(Y|X) = \sum_t y_t \log \lambda_t - \lambda_t$$

- log-likelihood is concave
- no local maxima [Paninski 04]



Now do it: tutorial2_spikehistcoupledGLM.m

Topics:

- 1-2. load data & bin spike trains
3. build design matrix, now including spike history!
4. fit Poisson GLM with spike-history
5. fit coupled Poisson GLM (4 neuron spike-history)
6. model comparison: log-likelihood & AIC

Block III: regularization

Modern statistics

- more dimensions than samples $D \geq N$

$$\begin{array}{c} \text{N} \\ \text{observations} \end{array} \left\{ \begin{bmatrix} y_1 \\ | \\ y_N \end{bmatrix} \right. = \begin{array}{c} \text{D regressors} \\ \overbrace{\begin{bmatrix} \text{---} \vec{x}_1 \text{---} \\ | \\ \text{---} \vec{x}_N \text{---} \end{bmatrix}} \end{array} \begin{bmatrix} w_1 \\ | \\ w_D \end{bmatrix} + \text{noise}$$



- fewer equations than unknowns!
- no unique solution

Simulated Example

- 100-element filter ($D=100$)
- 100 noisy samples ($N=100$)

true \mathbf{w}

maximum likelihood



maximize

$$\log p(\text{data}|\mathbf{w})$$


“overfitting” - parameters fit to details in the training data that are NOT useful for predicting new data

Simulated Example


- 100-element filter ($D=100$)
- 100 noisy samples ($N=100$)

$$\hat{\mathbf{w}} = (X^T X + \lambda I)^{-1} X^T Y$$

true \mathbf{w}




maximum likelihood



maximize

$$\log p(\text{data}|\mathbf{w})$$

“ridge regression”



maximize

$$\log p(\text{data}|\mathbf{w}) - \underbrace{\lambda \sum w_i^2}_{\text{penalty on big weights}}$$


penalty on
big weights

- biased, but gives improved performance for appropriate choice of λ (James & Stein 1960)


Simulated Example

- 100-element filter ($D=100$)
- 100 noisy samples ($N=100$)

true \mathbf{w}




maximum likelihood



maximize
 $\log p(\text{data}|\mathbf{w})$

“smoothed”



maximize
 $\log p(\text{data}|\mathbf{w})$
 $-\lambda \sum (w_i - w_{i-1})^2$
smoothness
penalty

Q: how to set the regularization strength λ ?
Simplest answer: use cross-validation!

Now do it: tutorial3_regularization.m

Topics:

- 1-2. load data, upsample to finer time bins.
3. divide into training and test data
4. maximum likelihood (linear-Gaussian model)
5. ridge regression (iid Gaussian prior on coeffs)
6. smoothing regression (iid Gaussian prior on diffs)